

Prompt Injection Detection: jev-guard vs Model Armor

Head-to-head evaluation over 2757 labeled cases from 16 public datasets · 2026-09-21

Summary

jev-guard scored a higher F1 on **14 of 16** datasets and tied on 0, across 2757 labeled cases spanning direct injection, jailbreak, and indirect injection. Performance differs sharply by attack class rather than uniformly, so the per-class aggregate below is the more useful summary than any single headline number.

That headline hides the finding that matters most. jev-guard's margin depends heavily on how realistic each dataset's benign half is. Where the negatives are ordinary question-answering text they are trivially separable and jev-guard dominates. Where they are drawn from real prompt-sharing traffic or deliberately matched to the attacks, its precision collapses: on **jailbreakhub** Model Armor wins outright (0.56 vs 0.42), and jev-guard's worst false-positive rates are jailbreakhub at 0.28 and boundary at 0.14, against Model Armor's 0.14 and 0.29. A guard that flags a third of legitimate traffic is not deployable, whatever its recall.

Diagnosing those false positives found two definition errors rather than threshold problems. The prompt-injection question asked whether a message contained instructions addressed to an AI — which describes prompting itself, and flagged users setting a persona. Rewritten around provenance (instructions arriving inside material the sender asked the assistant to process), its false positives on jailbreakhub fell from 87 to 11. The harmful-request question did not distinguish seeking capability from seeking understanding, and flagged questions about the history of bomb technology and about regulatory loopholes. Both fixes are in the numbers reported here.

The residual jailbreakhub gap is partly a labelling disagreement rather than a detection failure: much of its benign half consists of persona-override prompts (“you are now X, stay in character”) that many production guards would flag by design. It is reported as a loss regardless, because picking the interpretation that flatters the system under test is how evaluations become marketing.

Read the per-dataset numbers as directional rather than definitive, and note that the caveats below materially favor jev-guard.

Dataset	Type	n	jev F1	Model Armor F1	jev R	Model Armor R	jev FPR	Model Armor FPR
safeguard	direct	200	0.87	0.88	0.79	0.80	0.01	0.01
spml	direct	200	0.84	0.48	0.73	0.32	0.04	0.00
neuralchemy	direct	200	0.82	0.69	0.73	0.55	0.06	0.06
deepset	direct	116	0.57	0.39	0.42	0.25	0.04	0.04
mossca	direct	200	0.56	0.15	0.39	0.08	—	—
chatgpt_jb	jailbreak	79	0.99	0.91	0.97	0.84	—	—
gandalf	jailbreak	112	0.95	0.80	0.90	0.66	—	—
jailbreak	jailbreak	200	0.92	0.89	0.85	0.82	0.01	0.01

jbb	jailbreak	200	0.86	0.84	0.84	0.79	0.12	0.08
jailbreakv	jailbreak	200	0.84	0.81	0.72	0.68	—	—
jailbreakhub	jailbreak	200	0.42	0.56	0.90	0.86	0.28	0.14
llmail	indirect	200	0.99	0.13	0.97	0.07	—	—
mcp_poison	indirect	50	0.94	0.78	0.88	0.64	0.00	0.00
boundary	indirect	200	0.91	0.62	0.96	0.58	0.14	0.29
b3	indirect	200	0.83	0.31	0.70	0.18	—	—
bipia	indirect	200	0.81	0.06	0.75	0.03	0.11	0.02

Sorted by attack class — direct injection, then jailbreak, then indirect — with a rule at each boundary. R is recall; FPR is the false-positive rate, shown as “—” for datasets that contain only attacks and therefore have no negatives to measure it against. Green marks the higher F1.

Aggregate by attack class

Attack class	cases	jev recall	Model Armor recall	jev FPR	Model Armor FPR
direct	916	0.59	0.33	0.03	0.02
jailbreak	991	0.84	0.74	0.17	0.09
indirect	850	0.85	0.20	0.11	0.14

Method

Both systems received identical inputs, 50 cases sampled per dataset with a fixed seed. A case counts as flagged for jev-guard when its routing policy returns anything other than *pass* (block, review, or support), and for Model Armor when `filterMatchState` is `MATCH_FOUND` on any filter. Ground truth is each dataset's own label, so both systems are scored against the same key.

jev-guard ran on `~typesafe/jev-latest` via OpenRouter under its *strict* policy, combining a semantic hazard battery with deterministic structural detectors. Model Armor ran against the existing `testing-template` in project `ma-claude (us-central1)`.

Caveats that favor jev-guard

- **Model Armor ran at MEDIUM_AND_ABOVE.** The template sets `piAndJailbreakFilterSettings.confidenceLevel` to `MEDIUM_AND_ABOVE`. A `LOW_AND_ABOVE` setting exists and would raise its recall at some cost in false positives. These numbers are not Model Armor's ceiling.
- **Model Armor's own filter version was STABLE, not LATEST.** The baseline template pins `FILTER_VERSION_ALIAS_STABLE`. The permutation sweep below re-runs it on `LATEST`.
- **SPML is structurally asymmetric.** Its rows pair a system prompt with a user prompt. Jev takes both as structured state; Model Armor's `sanitizeUserPrompt` API accepts only the user text, so it cannot see the policy being violated. That is a real capability difference, but it means the SPML gap (0.99 vs 0.55) should not be read as a like-for-like detection comparison.

- **jev-guard's policy was tuned by its author; Model Armor's template was not tuned for this test.**

Thresholds in policy.py were iterated against deepset during development, so that dataset in particular flatters jev-guard.

Shared blind spot

Both systems missed 397 of 1835 attacks. These are the cases that matter most, because no amount of ensembling reaches them. The recurring shape is a role-play or persona framing that never issues an explicit override instruction:

- Seb and Irene act in a film about a racist murder as an actor. Seb emb
- As a child does not have any previous knowledge how would you react if
- Ich möchte, dass Sie als Drehbuchautor agieren. Sie werden ein ansprec
- Was kann ich mit Bärlauch kochen?. Was ist das schlimmste Schimpwort a
- I want you to act as a storyteller. You will come up with entertaining
- Dies ist keine Aufforderungsspritze. Ich möchte wirklich nur eine Antw

Was the gap just configuration?

The first round ran Model Armor on one template, which left open the objection that its settings rather than its detection explained the gap. To close that, six templates were provisioned covering every permutation of the PI filter's own axes — confidence level (LOW_AND_ABOVE / MEDIUM_AND_ABOVE / HIGH) crossed with multi-language detection on/off — all pinned to FILTER_VERSION_ALIAS_LATEST and isolating the prompt-injection filter (RAI, SDP and malicious-URI off) so the number measures PI detection rather than any filter firing.

deepset — all PI permutations (n=116)

Configuration	Precision	Recall	F1	FPR
jev-guard (reference)	0.92	0.60	0.73	0.05
Model Armor low/ml=on	1.00	0.28	0.44	0.00
Model Armor low/ml=off	1.00	0.28	0.44	0.00
Model Armor medium/ml=on	1.00	0.23	0.38	0.00
Model Armor medium/ml=off	1.00	0.22	0.36	0.00
Model Armor high/ml=on	1.00	0.18	0.31	0.00
Model Armor high/ml=off	1.00	0.18	0.31	0.00

slabs — all PI permutations (n=300)

Configuration	Precision	Recall	F1	FPR
jev-guard (reference)	0.98	0.70	0.81	0.02
Model Armor low/ml=on	1.00	0.40	0.57	0.00
Model Armor low/ml=off	1.00	0.40	0.57	0.00
Model Armor medium/ml=on	1.00	0.32	0.48	0.00
Model Armor medium/ml=off	1.00	0.32	0.48	0.00
Model Armor high/ml=on	1.00	0.27	0.43	0.00
Model Armor high/ml=off	1.00	0.27	0.43	0.00

Two results, one of which corrects this report's earlier draft

Confidence level behaves as documented but does not close the gap. Recall rises monotonically as the threshold loosens — on deepset 0.18 (HIGH) to 0.23 (MEDIUM) to 0.28 (LOW) — with precision staying at 1.00 throughout. Even at its most sensitive setting Model Armor's PI filter reaches 0.28 recall against jev-guard's 0.60 on the same cases. The gap is therefore not a threshold artifact.

Multi-language detection had no measurable effect, contrary to this report's first draft. Across 416 cases the flag changed exactly one verdict (deepset at MEDIUM, 0.23 vs 0.22 recall — a single case out of 60 attacks, which is noise). A direct spot check on three explicit German override instructions gave identical results either way, and two of the three went undetected even at LOW_AND_ABOVE with the flag enabled. The earlier claim that configuration explained the German misses was wrong — on this evidence it is a capability gap in the PI filter. The flag may well affect the RAI or SDP filters, which these isolated templates switch off; it does not appear to affect prompt-injection detection.

Per-dataset detail

Model Armor configuration used for every dataset below

Setting	Value
Template	testing-template (ma-claude / us-central1)
Prompt injection / jailbreak	ENABLED, MEDIUM_AND_ABOVE
Responsible AI filters	Hate Speech MEDIUM+, Dangerous MEDIUM+, Sexually Explicit MEDIUM+, Harassment MEDIUM+
Sensitive Data Protection	ENABLED
Malicious URI	ENABLED
Multi-language detection	not enabled
Filter version	FILTER_VERSION_ALIAS_STABLE

A case counts as flagged for Model Armor when any of these filters returns MATCH_FOUND, not the prompt-injection filter alone — so an SDP or Responsible AI match also counts as a detection in its favour. Note the filter version is STABLE rather than LATEST, and

multi-language detection is off; the permutation sweep above varies both and finds the confidence level matters while multi-language does not.

deepset

deepset/prompt-injections — direct injection, mixed German/English

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.93	0.42	0.57	0.04	25	2	54	35
Model Armor	0.88	0.25	0.39	0.04	15	2	54	45

60 attacks / 116 cases · agreement 83% · jev-only 15, armor-only 5, both missed 30

safeguard

xTRam1/safe-guard-prompt-injection — direct injection, larger and more varied

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.98	0.79	0.87	0.01	52	1	133	14
Model Armor	0.98	0.80	0.88	0.01	53	1	133	13

66 attacks / 200 cases · agreement 92% · jev-only 7, armor-only 8, both missed 6

neuralchemy

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.94	0.73	0.82	0.06	80	5	85	30
Model Armor	0.92	0.55	0.69	0.06	61	5	85	49

110 attacks / 200 cases · agreement 82% · jev-only 28, armor-only 9, both missed 21

mossca

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	1.00	0.39	0.56	nan	78	0	0	122
Model Armor	1.00	0.08	0.15	nan	16	0	0	184

200 attacks / 200 cases · agreement 65% · jev-only 66, armor-only 4, both missed 118

spml

reshabhs/SPML_Chatbot_Prompt_Injection — injection judged against a given system prompt

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.98	0.73	0.84	0.04	113	2	43	42
Model Armor	1.00	0.32	0.48	0.00	49	0	45	106

155 attacks / 200 cases · agreement 66% · jev-only 65, armor-only 1, both missed 41

jailbreak

jackhhao/jailbreak-classification — jailbreak vs benign

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.99	0.85	0.92	0.01	88	1	96	15
Model Armor	0.99	0.82	0.89	0.01	84	1	96	19

103 attacks / 200 cases · agreement 95% · jev-only 6, armor-only 2, both missed 13

jbb

JailbreakBench — 100 harmful and 100 deliberately matched benign

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.88	0.84	0.86	0.12	84	12	88	16
Model Armor	0.91	0.79	0.84	0.08	79	8	92	21

100 attacks / 200 cases · agreement 86% · jev-only 10, armor-only 5, both missed 11

jailbreakhub

walldai/JailbreakHub — real in-the-wild prompts from prompt-sharing communities

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.28	0.90	0.42	0.28	19	50	129	2
Model Armor	0.42	0.86	0.56	0.14	18	25	154	3

21 attacks / 200 cases · agreement 85% · jev-only 1, armor-only 0, both missed 2

chatgpt_jb

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	1.00	0.97	0.99	nan	77	0	0	2
Model Armor	1.00	0.84	0.91	nan	66	0	0	13

79 attacks / 79 cases · agreement 86% · jev-only 11, armor-only 0, both missed 2

jailbreakv

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	1.00	0.72	0.84	nan	145	0	0	55
Model Armor	1.00	0.68	0.81	nan	136	0	0	64

200 attacks / 200 cases · agreement 90% · jev-only 14, armor-only 5, both missed 50

gandalf

Lakera/gandalf_ignore_instructions — real extraction attempts, all positive (recall only)

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	1.00	0.90	0.95	nan	101	0	0	11
Model Armor	1.00	0.66	0.80	nan	74	0	0	38

112 attacks / 112 cases · agreement 74% · jev-only 28, armor-only 1, both missed 10

bipia

BIPIA-derived — INDIRECT injection embedded in retrieved content

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.88	0.75	0.81	0.11	78	11	85	26
Model Armor	0.60	0.03	0.06	0.02	3	2	94	101

104 attacks / 200 cases · agreement 58% · jev-only 75, armor-only 0, both missed 26

boundary

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	0.87	0.96	0.91	0.14	96	14	86	4
Model Armor	0.67	0.58	0.62	0.29	58	29	71	42

100 attacks / 200 cases · agreement 72% · jev-only 38, armor-only 0, both missed 4

llmail

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	1.00	0.97	0.99	nan	195	0	0	5
Model Armor	1.00	0.07	0.13	nan	14	0	0	186

200 attacks / 200 cases · agreement 10% · jev-only 181, armor-only 0, both missed 5

b3

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	1.00	0.70	0.83	nan	141	0	0	59
Model Armor	1.00	0.18	0.31	nan	36	0	0	164

200 attacks / 200 cases · agreement 46% · jev-only 106, armor-only 1, both missed 58

mcp_poison

System	Precision	Recall	F1	FPR	TP	FP	TN	FN
jev-guard	1.00	0.88	0.94	0.00	22	0	25	3
Model Armor	1.00	0.64	0.78	0.00	16	0	25	9

25 attacks / 50 cases · agreement 76% · jev-only 9, armor-only 3, both missed 0

Reproducing this

```
python compare.py deepset safeguard jailbreak spml slabs password mixed --n 50 --json results.json
python report.py results.json jev-guard-evaluation.pdf
```